

p8105_hw5_ruijipan

ruijipan

2022-11-13

Introduction

This report is used to explain assignment 5 of R language. Assignment 5 mainly focuses on the study and training of iteration grammar in R language.

Problem 1

First, create a data folder to store the data needed for this report; Under the data folder, create a subfolder named longitudinal-study to store corresponding dataset.

The following code shows how to use the `list.files` function and the `map_dfr` function to read data sets in batches and store the data as a data frame type R language object; Among them, the `paste` function is used for string merging, which combines the path with the dataset. The file names are merged to form the input parameters of the `read.csv` function.

```
> library(tidyverse)
> filenames = list.files('./data/longitudinal-study')
> filepaths = paste('./data/longitudinal-study/',filenames,sep="")
> df = map_dfr(filepaths, read.csv, .id = "input")
> df
```

	input	week_1	week_2	week_3	week_4	week_5	week_6	week_7	week_8
1	1	0.20	-1.31	0.66	1.96	0.23	1.09	0.05	1.94
2	2	1.13	-0.88	1.07	0.17	-0.83	-0.31	1.58	0.44
3	3	1.77	3.11	2.22	3.26	3.31	0.89	1.88	1.01
4	4	1.04	3.66	1.22	2.33	1.47	2.70	1.87	1.66
5	5	0.47	-0.58	-0.09	-1.37	-0.32	-2.17	0.45	0.48
6	6	2.37	2.50	1.59	-0.16	2.08	3.07	0.78	2.35
7	7	0.03	1.21	1.13	0.64	0.49	-0.12	-0.07	0.46
8	8	-0.08	1.42	0.09	0.36	1.18	-1.16	0.33	-0.44
9	9	0.08	1.24	1.44	0.41	0.95	2.75	0.30	0.03
10	10	2.14	1.15	2.52	3.44	4.26	0.97	2.73	-0.53
11	11	3.05	3.67	4.84	5.80	6.33	5.46	6.38	5.91
12	12	-0.84	2.63	1.64	2.58	1.24	2.32	3.11	3.78
13	13	2.15	2.08	1.82	2.84	3.36	3.61	3.37	3.74
14	14	-0.62	2.54	3.78	2.73	4.49	5.82	6.00	6.49
15	15	0.70	3.33	5.34	5.57	6.90	6.66	6.24	6.95
16	16	3.73	4.08	5.40	6.41	4.87	6.09	7.66	5.83
17	17	1.18	2.35	1.23	1.17	2.02	1.61	3.13	4.88
18	18	1.37	1.43	1.84	3.60	3.80	4.72	4.68	5.70
19	19	-0.40	1.08	2.66	2.70	2.80	2.64	3.51	3.27
20	20	1.09	2.80	2.80	4.30	2.25	6.57	6.09	4.64

Next, clean the dataset. As you can see, after the data is read in, it is a data frame type object. Rename the input field to arm by using the `rename` function, and use the character in front of the symbol “ ” in the file

name as the data content of this field; Then, a new variable "subject_id" is created, and the character after the symbol "" in the file name is used as the data content of this field. After that, use the select function to rearrange the data fields. As shown in the table below, the data is relatively clean, so no further data cleaning is required.

```
> df = rename(df, c(arm=input))
> str_arm = strsplit(filenamees, "\\.")
>
> var_arm = vector("list", length = 20)
> var_id = vector("list", length = 20)
> for(i in 1:20){
+   var_arm[[i]] = strsplit(str_arm[[i]][1], "_")[[1]][1]
+   var_id[[i]] = strsplit(str_arm[[i]][1], "_")[[1]][2]
+ }
> df$arm = var_arm
> df$subject_id = var_id
> df = df %>%
+   select(arm, subject_id, everything())
> df
```

	arm	subject_id	week_1	week_2	week_3	week_4	week_5	week_6	week_7	week_8
1	con	01	0.20	-1.31	0.66	1.96	0.23	1.09	0.05	1.94
2	con	02	1.13	-0.88	1.07	0.17	-0.83	-0.31	1.58	0.44
3	con	03	1.77	3.11	2.22	3.26	3.31	0.89	1.88	1.01
4	con	04	1.04	3.66	1.22	2.33	1.47	2.70	1.87	1.66
5	con	05	0.47	-0.58	-0.09	-1.37	-0.32	-2.17	0.45	0.48
6	con	06	2.37	2.50	1.59	-0.16	2.08	3.07	0.78	2.35
7	con	07	0.03	1.21	1.13	0.64	0.49	-0.12	-0.07	0.46
8	con	08	-0.08	1.42	0.09	0.36	1.18	-1.16	0.33	-0.44
9	con	09	0.08	1.24	1.44	0.41	0.95	2.75	0.30	0.03
10	con	10	2.14	1.15	2.52	3.44	4.26	0.97	2.73	-0.53
11	exp	01	3.05	3.67	4.84	5.80	6.33	5.46	6.38	5.91
12	exp	02	-0.84	2.63	1.64	2.58	1.24	2.32	3.11	3.78
13	exp	03	2.15	2.08	1.82	2.84	3.36	3.61	3.37	3.74
14	exp	04	-0.62	2.54	3.78	2.73	4.49	5.82	6.00	6.49
15	exp	05	0.70	3.33	5.34	5.57	6.90	6.66	6.24	6.95
16	exp	06	3.73	4.08	5.40	6.41	4.87	6.09	7.66	5.83
17	exp	07	1.18	2.35	1.23	1.17	2.02	1.61	3.13	4.88
18	exp	08	1.37	1.43	1.84	3.60	3.80	4.72	4.68	5.70
19	exp	09	-0.40	1.08	2.66	2.70	2.80	2.64	3.51	3.27
20	exp	10	1.09	2.80	2.80	4.30	2.25	6.57	6.09	4.64

Problem 2

The problem 2 uses the homicides data collected by The Washington Post. This data set contains a total of 52,179 observations and 12 variables. It mainly includes the ID of the case, the reported date of the case, some basic information of the victim, the city where the case happened and the settlement of the case. The following code creates a new city_state field for this data set, merges the city field and the state field, and calculates the total number of cases and the number of unsolved cases in each city.

```
> homicides_df = read.csv("../data/homicides/homicide-data.csv")
> homicides_df$city_state = paste(homicides_df$city, homicides_df$state)
> head(homicides_df)
```

	uid	reported_date	victim_last	victim_first	victim_race	victim_age
1	Alb-000001	20100504	GARCIA	JUAN	Hispanic	78
2	Alb-000002	20100216	MONTOYA	CAMERON	Hispanic	17

```

3 Alb-000003      20100601 SATTERFIELD      VIVIANA      White      15
4 Alb-000004      20100101      MENDIOLA      CARLOS      Hispanic      32
5 Alb-000005      20100102      MULA      VIVIAN      White      72
6 Alb-000006      20100126      BOOK      GERALDINE      White      91
  victim_sex      city state      lat      lon      disposition
1      Male Albuquerque      NM 35.09579 -106.5385549 Closed without arrest
2      Male Albuquerque      NM 35.05681 -106.715321      Closed by arrest
3      Female Albuquerque      NM 35.08609 -106.695568 Closed without arrest
4      Male Albuquerque      NM 35.07849 -106.5560938 Closed by arrest
5      Female Albuquerque      NM 35.13036 -106.5809862 Closed without arrest
6      Female Albuquerque      NM 35.15111 -106.537797      Open/No arrest
  city_state
1 Albuquerque NM
2 Albuquerque NM
3 Albuquerque NM
4 Albuquerque NM
5 Albuquerque NM
6 Albuquerque NM

```

```

> homicidesBycity =
+   homicides_df %>%
+   group_by(city) %>%
+   summarise(count = n())
> head(homicidesBycity)
# A tibble: 6 x 2
  city      count
  <chr>    <int>
1 Albuquerque  378
2 Atlanta      973
3 Baltimore   2827
4 Baton Rouge  424
5 Birmingham   800
6 Boston      614

```

Next, the function `prop.test` is used to estimate the unsolved proportion of cases in “Baltimore, MD” city (about 65%), and the list is visualized to output the estimated proportion and confidence intervals of the city.

```

> unsolved_homicides =
+   homicides_df %>%
+   filter(disposition %in% c("Closed without arrest", "Open/No arrest")) %>%
+   summarise(disposition = "unsolved", total = n())
> head(unsolved_homicides)
  disposition total
1      unsolved 26505

```

```

> Baltimore_MD_unsolved =
+   homicides_df %>%
+   filter(disposition %in% c("Closed without arrest", "Open/No arrest") & city_state == "Baltimore MD")
+   summarise(disposition = "MD_unsolved", total = n())
> head(Baltimore_MD_unsolved)
  disposition total
1 MD_unsolved  1825

```

```

> Baltimore_MD_total =
+   homicides_df %>%
+   filter(city_state == "Baltimore MD") %>%

```

```

+ summarise(disposition = "MD_total",total = n())
> head(Baltimore_MD_total)
  disposition total
1    MD_total  2827

> prop_test = prop.test(Baltimore_MD_unsolved$total,Baltimore_MD_total$total)
> prop_test_df = broom::tidy(prop_test)
> prop_test_result_df = prop_test_df %>%
+   select(estimate, conf.low, conf.high)
> prop_test_result_df
# A tibble: 1 x 3
  estimate conf.low conf.high
    <dbl>    <dbl>    <dbl>
1   0.646   0.628   0.663

> cat("proportion estimate: ",prop_test_df$estimate, "\n")
proportion estimate:  0.6455607
> cat("the 0.95 conf.low: ", prop_test_df$conf.low, "\n")
the 0.95 conf.low:  0.6275625
> cat("the 0.95 conf.high: ", prop_test_df$conf.high, "\n")
the 0.95 conf.high:  0.6631599

```

For the above process, we package it into a function, and input it into different cities to get different results.

```

> # define a function for the above code
> # input: city
> # output: prop_test_result_df
> proportion = function(x){
+   city_total =
+   homicides_df %>%
+   filter(city_state == x) %>%
+   summarise(disposition = "city_total",total = n())
+
+   city_unsolved =
+   homicides_df %>%
+   filter(disposition %in% c("Closed without arrest","Open/No arrest") & city_state == x) %>%
+   summarise(disposition = "city_unsolved",total = n())
+
+   prop_test = prop.test(city_unsolved$total,city_total$total)
+   prop_test_df = broom::tidy(prop_test)
+   prop_test_result_df = prop_test_df %>%
+     select(estimate, conf.low, conf.high)
+
+   prop_test_result_df
+ }

```

After that, use the `map_dfr` function to run our packaged functions in batches and output a data frame containing the output results of all cities.

```

>
> prop_result = map_dfr(unique(homicides_df$city_state), proportion, .id = "input")
> prop_result = rename(prop_result,c(city_state=input))
> prop_result$city_state = unique(homicides_df$city_state)
> prop_result
# A tibble: 52 x 4
  city_state estimate conf.low conf.high

```

```

  <chr>          <dbl>    <dbl>    <dbl>
1 Albuquerque NM 0.386    0.337    0.438
2 Atlanta GA     0.383    0.353    0.415
3 Baltimore MD   0.646    0.628    0.663
4 Baton Rouge LA 0.462    0.414    0.511
5 Birmingham AL  0.434    0.399    0.469
6 Boston MA      0.505    0.465    0.545
7 Buffalo NY     0.612    0.569    0.654
8 Charlotte NC   0.300    0.266    0.336
9 Chicago IL     0.736    0.724    0.747
10 Cincinnati OH 0.445    0.408    0.483
# ... with 42 more rows


```

For each city, use ggplot to visualize the proportion of its unsolved cases, and its corresponding confidence interval.

```

> prop_result %>%
+   ggplot() +
+   geom_errorbar(aes(x=city_state, ymin=conf.low, ymax=conf.high,color=city_state), position = position_dodge())
+   ggtitle("unsolved proportion for city_state")

```



Problem 3

For problem 3, use the `set.seed` function to ensure that our results can be reproduced; Then, the data with mean value of 0 and standard deviation of 5 are generated by cyclic batch. The total number of iterations is 5,000, and each iteration generates 30 observation data.

```

> library(tidyverse)
> set.seed(1)

```

```

> data_norm = vector("list", 5000)
> for(i in 1:5000){
+   data_norm[[i]] = rnorm(n = 30, mean = 0, sd = 5)
+ }
> listcol_df =
+   tibble(
+     sample_id = c(1:5000),
+     samp = data_norm
+   )
>
> listcol_df =
+   listcol_df %>%
+   mutate(summary = map2(.x = samp, .y = 0, ~t.test(x = .x, mu = .y)))

```

For the above-mentioned generated data, 5000 datasets are tested by single sample mean, and the estimated value and significance are extracted. The confidence level of 0.05 is used to judge whether the test result of each data set is significant, and whether the original hypothesis is rejected or not is stored as a rejected variable.

```

> mean_test_result = map_dfr(listcol_df[[3]], broom::tidy, .id = "sample_id") %>%
+   select(sample_id, estimate, p.value) %>%
+   mutate(rejected = p.value > 0.05)
> mean_test_result
# A tibble: 5,000 x 4
  sample_id estimate p.value rejected
  <chr>         <dbl>   <dbl> <lgl>
1 1          0.412  0.629 TRUE
2 2          0.664  0.368 TRUE
3 3          0.551  0.534 TRUE
4 4          0.567  0.487 TRUE
5 5         -1.65   0.0599 TRUE
6 6          1.19   0.229 TRUE
7 7          0.334  0.738 TRUE
8 8         -1.19   0.209 TRUE
9 9          0.122  0.887 TRUE
10 10         0.684  0.472 TRUE
# ... with 4,990 more rows

```

The above process is packaged into a function, the input parameter of the function is the average value of random numbers with normal distribution, and the output is the test result of each data set.

```

> # define a function for the above code
> # input: mean
> # output: dataframe containing sample_id, true_u, estimate, p.value, rejected
> library(tidyverse)
> mean_test = function(x){
+   set.seed(1)
+   data_norm = vector("list", 5000)
+   for(i in 1:5000){
+     data_norm[[i]] = rnorm(n = 30, mean = x, sd = 5)
+   }
+   listcol_df =
+     tibble(
+       sample_id = c(1:5000),
+       samp = data_norm
+     )
+ }

```

```

+   )
+
+   listcol_df =
+     listcol_df %>%
+     mutate(summary = map2(.x = samp, .y = x, ~t.test(x = .x, mu = .y)))
+
+   mean_test_result = map_dfr(listcol_df[[3]], broom::tidy, .id = "sample_id") %>%
+   select(sample_id, estimate, p.value) %>%
+   mutate(true_u = x, rejected = p.value > 0.05)
+   mean_test_result
+ }

```

Take the mean values of 1, 2, 3, 4, 5, 6, respectively, and run the above functions in batches. The results of the generated functions are stored in the variable `mean_test_result_df`.

```

> mean_vec = c(1,2,3,4,5,6)
> mean_test_result_df = map_dfr(mean_vec, mean_test)

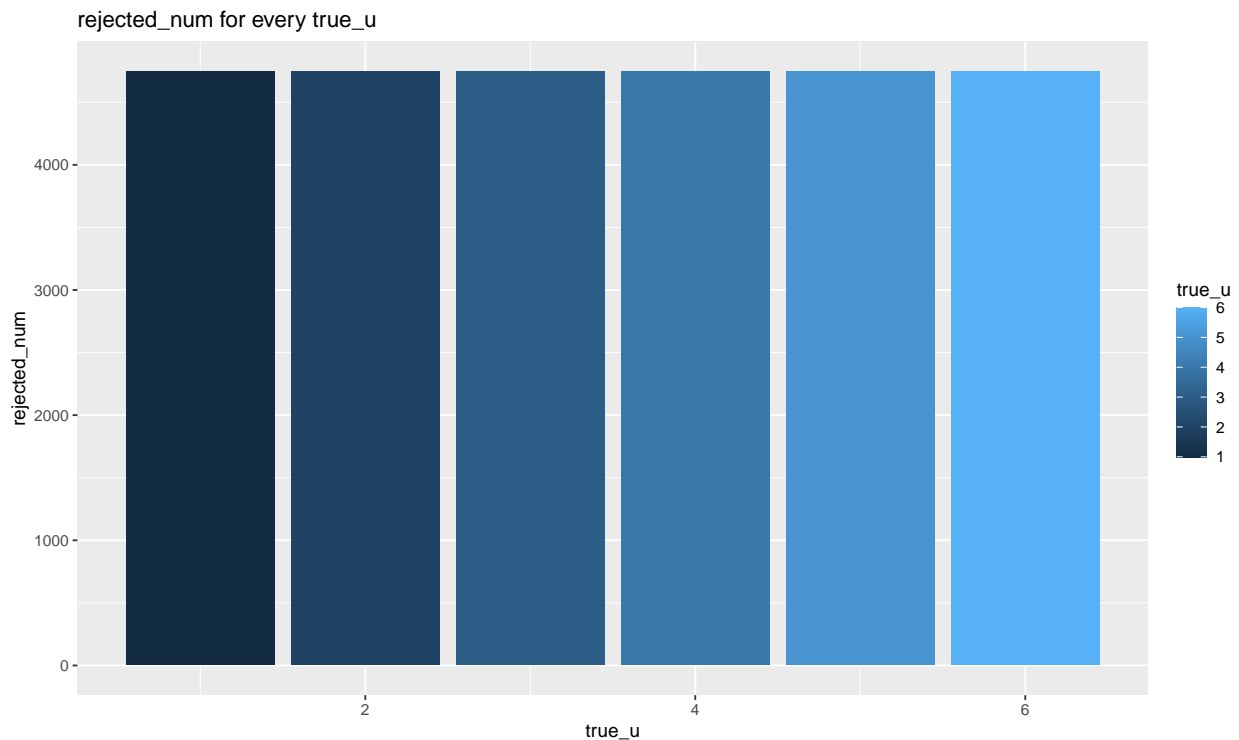
```

For the test result list generated above, `ggplot` is used to show the number of data sets that fail the test under each different mean value. As shown below:

```

> mean_test_result_df %>%
+   group_by(true_u) %>%
+   summarise(rejected_num = sum(rejected)) %>%
+   ggplot(aes(x = true_u, y = rejected_num, fill = true_u)) +
+   geom_bar(stat = "identity") +
+   ggtitle("rejected_num for every true_u")

```

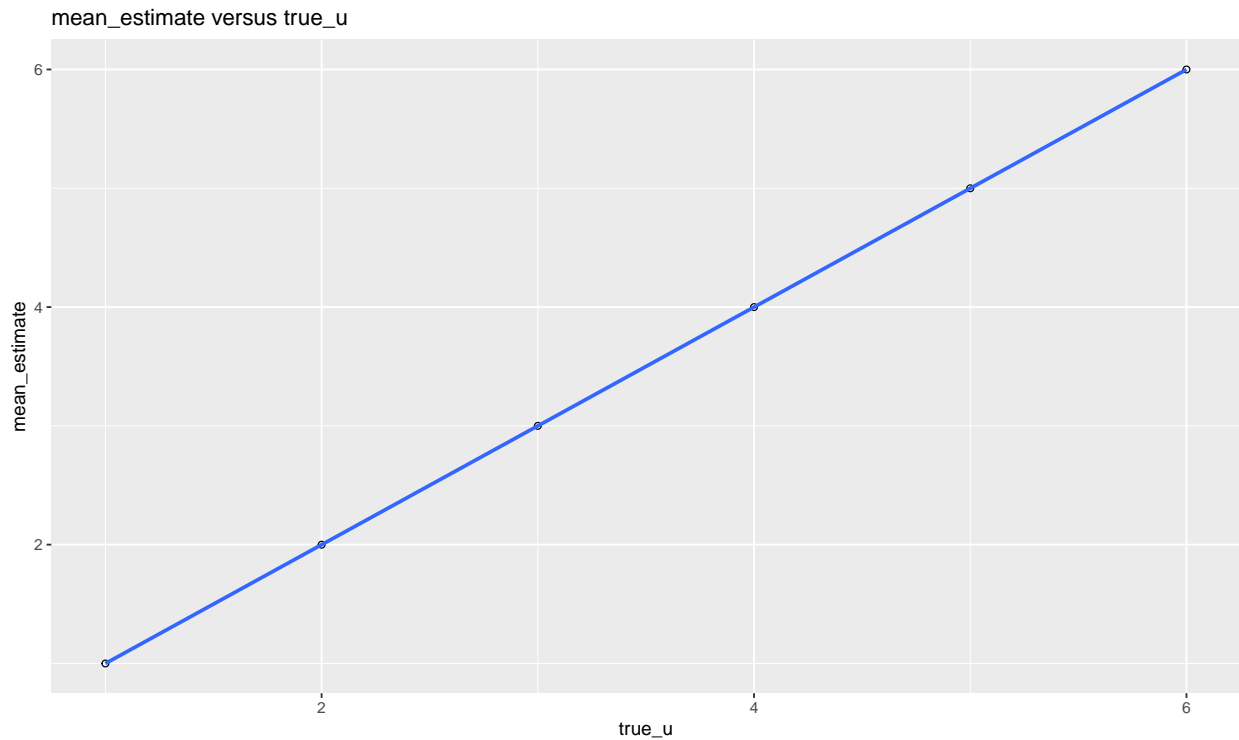


Similarly, the list of test results is grouped according to `true_u`, and the estimated mean value is calculated, and the corresponding scatter plot is drawn.

```

> mean_test_result_df %>%
+   group_by(true_u) %>%
+   summarise(mean_estimate = mean(estimate)) %>%
+   ggplot(aes(x = true_u, y = mean_estimate)) +
+   geom_point(shape=1) +
+   geom_smooth(method = 'loess') +
+   ggtitle("mean_estimate versus true_u")

```

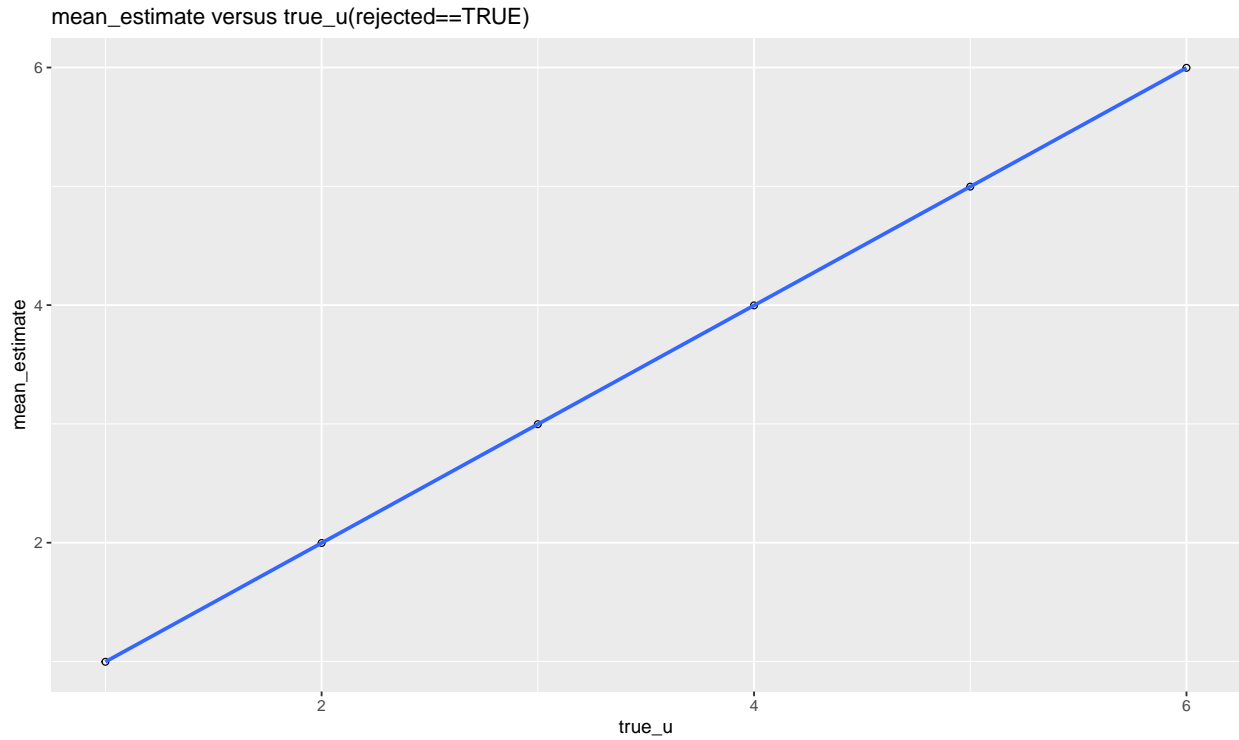


Group the list of test results according to true_u, first screen out the data that reject the original hypothesis, then calculate the estimated mean and draw the corresponding scatter plot.

```

> mean_test_result_df %>%
+   filter(rejected == TRUE) %>%
+   group_by(true_u) %>%
+   summarise(mean_estimate = mean(estimate)) %>%
+   ggplot(aes(x = true_u, y = mean_estimate)) +
+   geom_point(shape=1) +
+   geom_smooth(method = 'loess') +
+   ggtitle("mean_estimate versus true_u(rejected==TRUE)")

```

It can be found that the estimated mean values calculated by the two scatterplots are almost identical. This is because the percentage of failed tests is very high, reaching 80%~90%. Therefore, the estimated average is of course mainly determined by the data sets that failed tests.